



Designing CNN for Optical Character Recognition

Sandesh Kumar-1973056
Faisal Gul - 1950572

What is Optical Character Recognition (OCR)



Popular challenge in Computer Vision



Optical Character Recognition (OCR) is the process of detecting and reading text in images



Handwritten OCR can work with characters captured at various angles and shapes

AIM of the Project

- Build an OCR system by implementing a CNN architecture to distinguish between handwritten letters from the EMNIST Letters dataset.
- Evaluate CNN model on handwritten letters captured by camera.

What is EMNIST Dataset



EMNIST dataset is a set of handwritten character digits derived from the NIST Special Database 19 and converted to a 28x28 pixel image format.



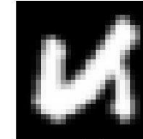
It consists of a set of upper- and lowercase letters, categorized into 26 classes.



The dataset is balanced.



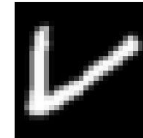
31 (31)



35 (35)



30 (30)



7 (7)



15 (15)



21 (21)



19 (19)



42 (42)



36 (36)

Test datasets

1

6.000 images from
EMNIST Letters



(a) Test 1

2

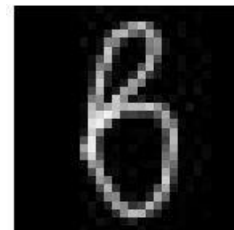
52 Handwritten
letters captured by
camera



(b) Test 2

3

52 letters drawn with
the use of Tkinter user
interface



(c) Test 3

Loading EMNIST data

Vast amount of data in EMNIST



Download from NIST site as a zipped MATLAB file



Temporarily whole dataset stored



Needed data pulled out, rest out of memory

```
def load_data(path="emnist_matlab.npz", type="letters"):
```

```
    # Log about loading
```

```
    logging.basicConfig(level=logging.INFO)
```

```
    logging.info('Loading dataset = emnist')
```

```
    # Load data
```

```
    path = get_file(
```

```
        path, origin=("http://www.itl.nist.gov/iaui/vip/cs\_links/EMNIST/"  
                      "matlab.zip")
```

```
)
```

```
    with ZipFile(path, "r") as opened_zip:
```

```
        # Read file and temporarily store it
```

```
        file_name = f"./{type}.mat"
```

```
        f = open(file_name, "wb")
```

```
        f.write(opened_zip.read(f"matlab/emnist-{type}.mat"))
```

```
        f.close()
```

```
    # Load data from Matlab file.
```

```
    # Source: https://stackoverflow.com/a/53547262
```

```
    mat = sio.loadmat(file_name)
```

```
    data = mat["dataset"]
```

```
    input_train = data["train"][0, 0]["images"][0, 0]
```

```
    target_train = data["train"][0, 0]["labels"][0, 0].flatten()
```

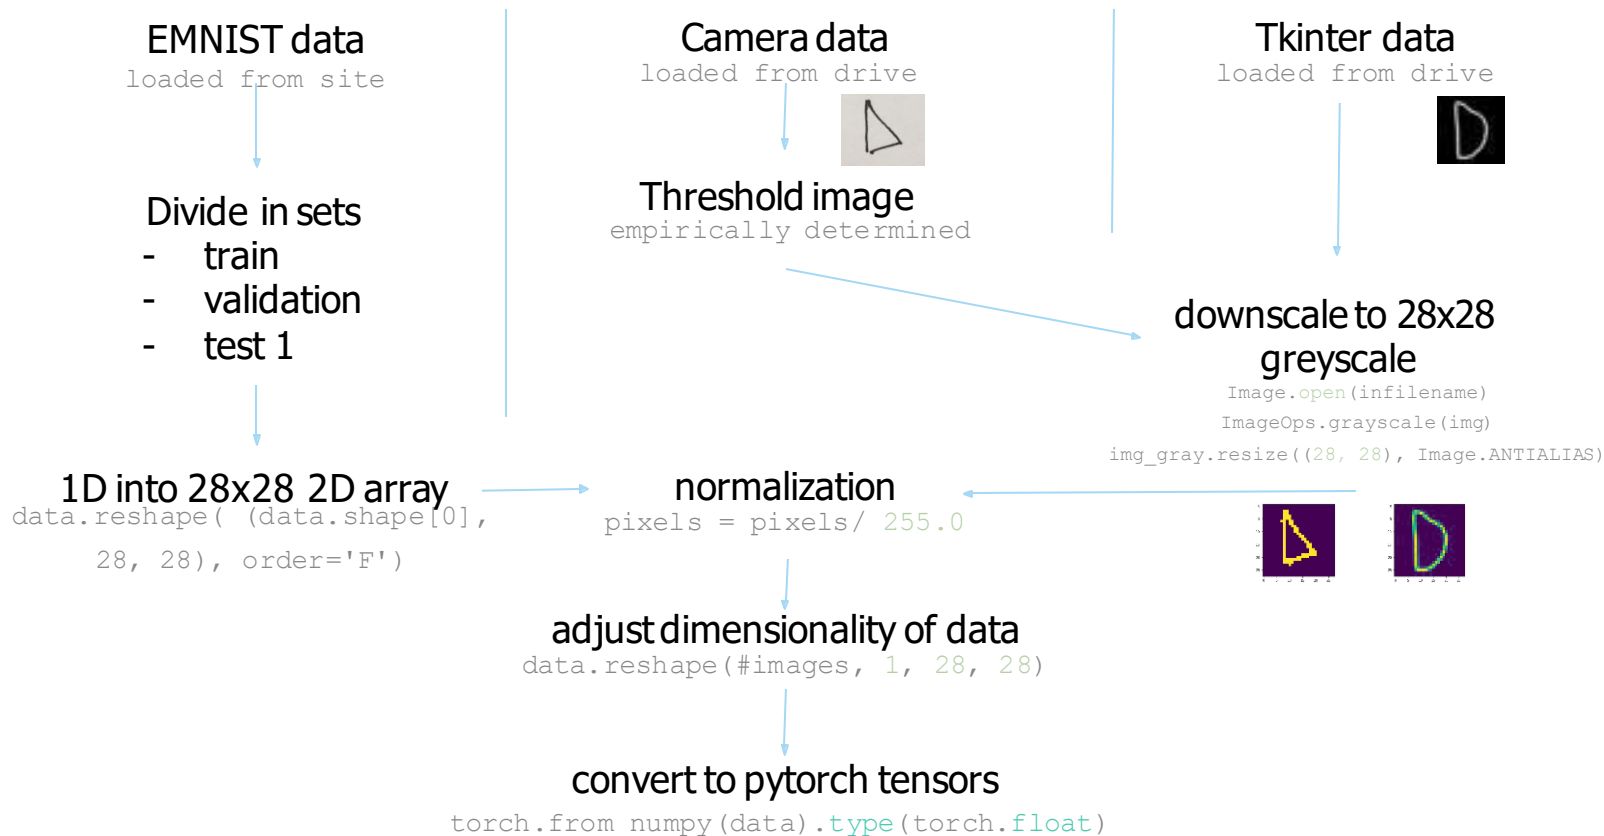
```
    input_test = data["test"][0, 0]["images"][0, 0]
```

```
    target_test = data["test"][0, 0]["labels"][0, 0].flatten()
```

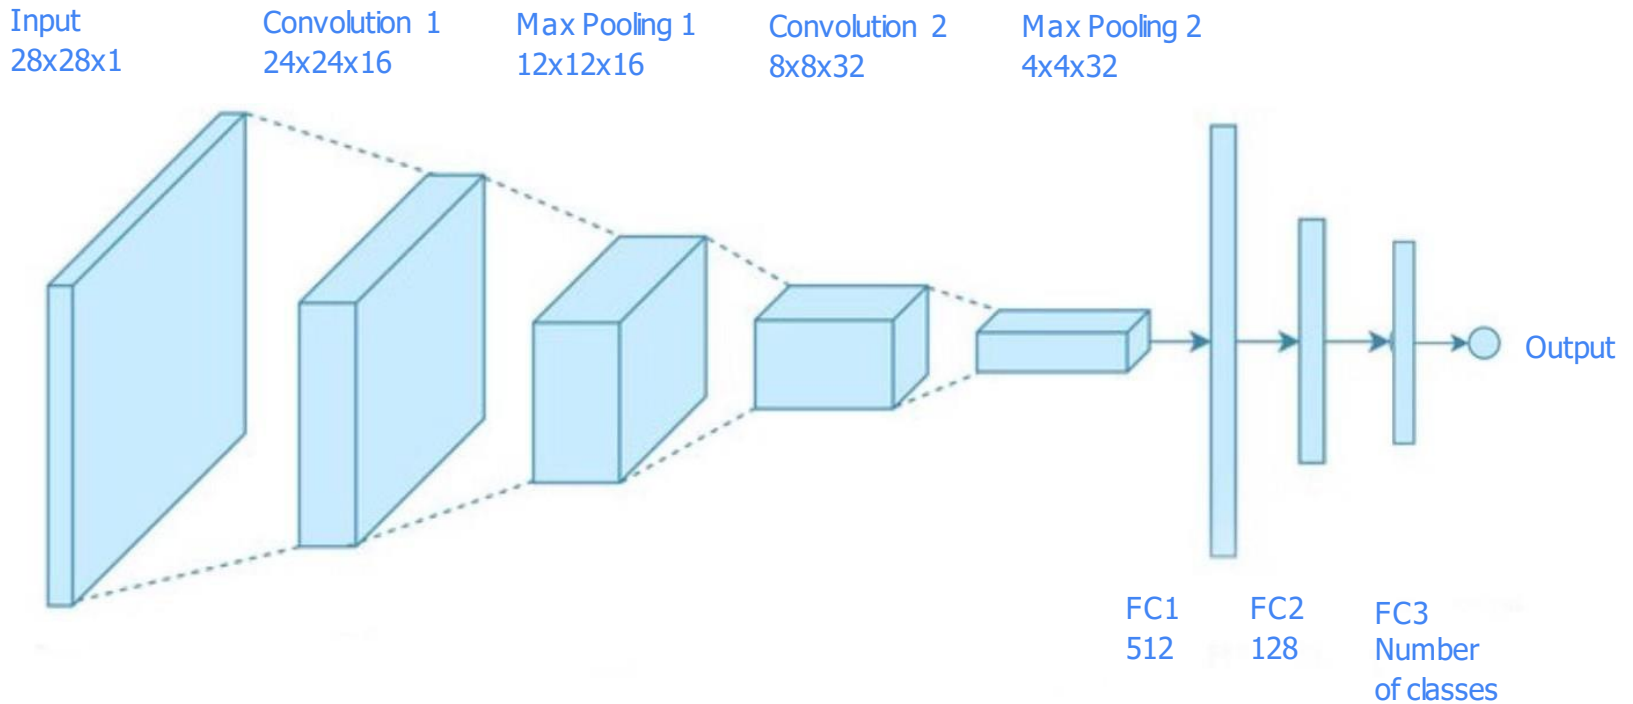
```
    # Remove data when loaded
```

```
    os.remove(file_name)
```

Preprocessing image data



Model architecture




```
class CNN(nn.Module):
```

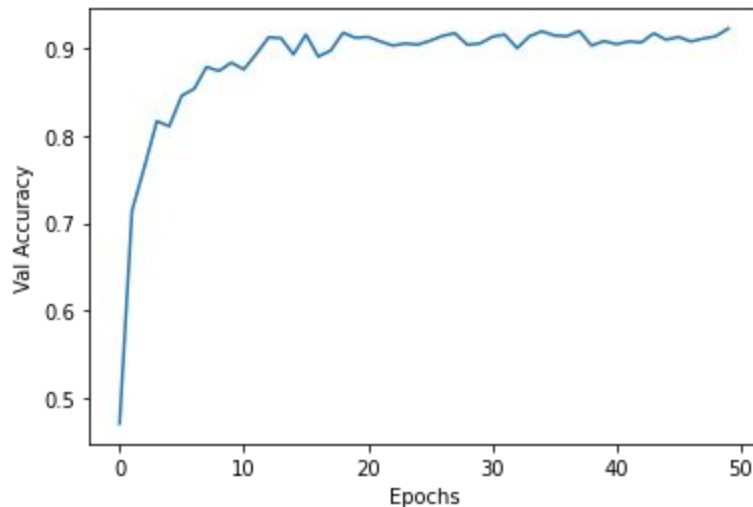
```
    def __init__(self, num_classes=27):  
        super(CNN, self).__init__()  
        self.conv = nn.Conv2d(1, 16, 5)  
        self.pool = nn.MaxPool2d(2)  
        self.conv2 = nn.Conv2d(16, 32, 5)  
        self.drop = nn.Dropout(p=0.2)  
        self.pool2 = nn.MaxPool2d(2)  
        self.fc1 = nn.Linear(32*4*4, 512)  
        self.fc2 = nn.Linear(512, 128)  
        self.fc3 = nn.Linear(128, num_classes)
```

```
    def forward(self, x):  
        x = torch.nn.functional.relu(self.conv(x))  
        x = self.pool(x)  
        x = torch.nn.functional.relu(self.conv2(x))  
        x = self.drop(x)  
        x = self.pool2(x)  
        x = torch.flatten(x, 1)  
        x = torch.nn.functional.relu(self.fc1(x))  
        x = torch.nn.functional.relu(self.fc2(x))  
        x = self.fc3(x)
```

```
    return x
```

```
    cnn = CNN()
```

Model architecture



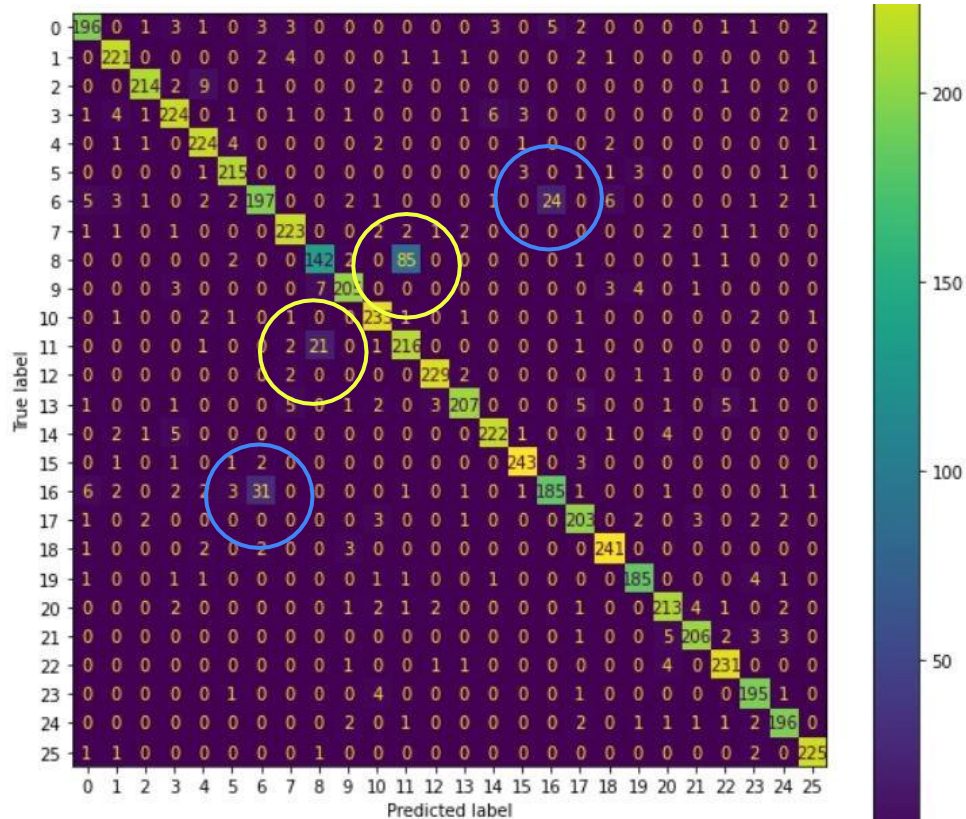
Evaluation of the model using test statistics

Botlab 99.2%
Peng and Yin 88.77%-99.75%



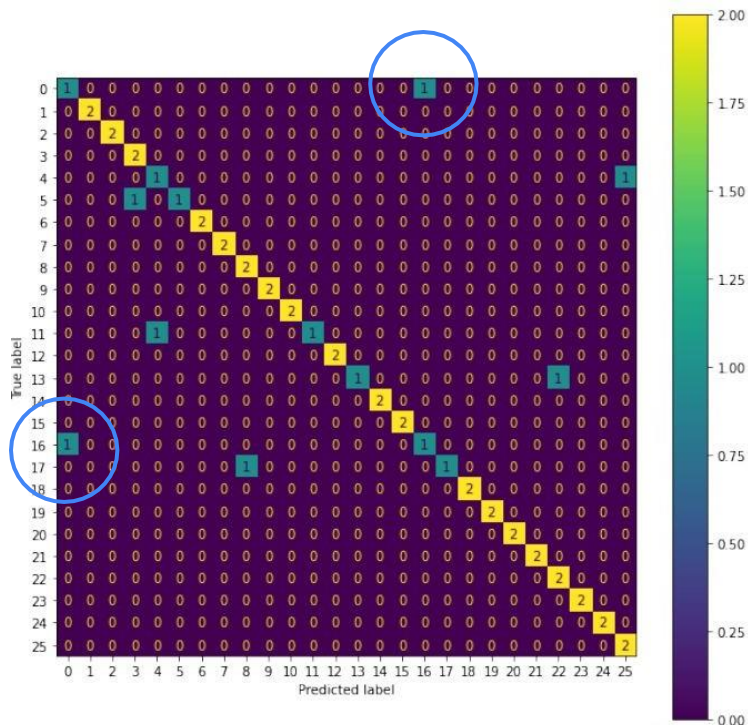
Test set number & sample size	Accuracy	F1 score	Precision score	Recall score
1 (6000) EMNIST letters	0.92	0.91	0.92	0.92
2 (52) Camera captured	0.87	0.86	0.89	0.87
3 (52) Tk UI	0.88	0.88	0.92	0.88

Confusion Matrix Test 1

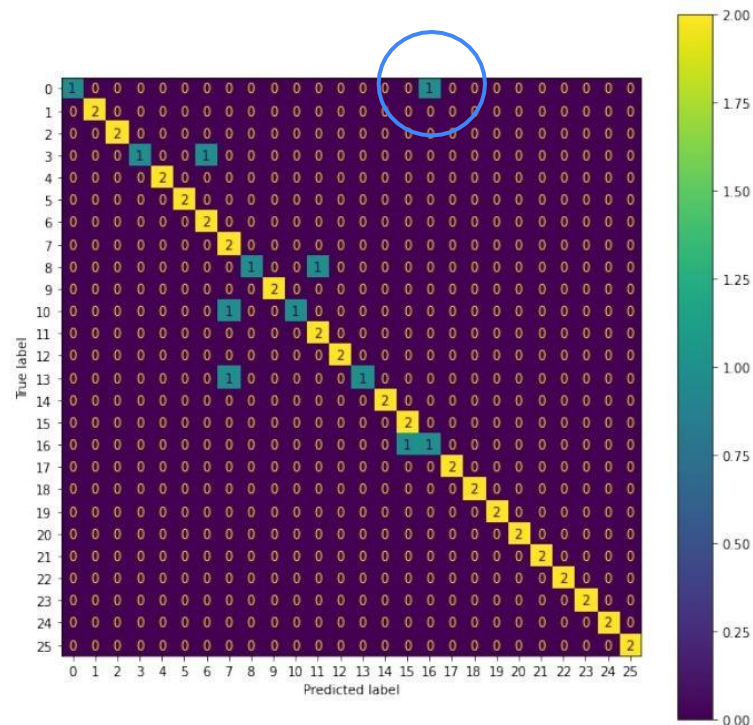


# of mislabeled cases	Letter 1	Letter 2
106	I	L
55	G	Q

Confusion Matrix Test 2 and Test 3



(a) Test 1



(b) Test 2

Mislabeling of A and Q letters together
having 3 wrong predictions out of 8

Limitations and Challenges



The EMNIST Letters dataset is too big, therefore, we used only its part. But we still achieved validation accuracy of 92 %



Test2 images were captured by camera during daylight. But when we ran model on more dark images, model performance was poor



The model generally does not show good performance on some images, because they are very different from EMNIST data

A decorative graphic on the left side of the slide consisting of a grid of hexagons. Some hexagons are filled with light blue or light green, while others are empty, creating a pattern that resembles a honeycomb or a molecular structure. The hexagons are outlined in a darker blue or green color.

Conclusion

Automatic recognition of handwritten letters, captured by camera, can be implemented using the Pytorch CNN framework and showed an accuracy of 87%